

# **Achieving High Quality of Attribution in Provenance-based Intrusion Detection Systems**

---

Tristan Bilot, University of British Columbia  
2025/12/05

# What is this talk about?

Mainly on 2 papers published at **USENIX Sec'25**.

- *ORTHRUS: Achieving High Quality of Attribution in Provenance-based Intrusion Detection Systems*
- *Sometimes Simpler is Better: A Comprehensive Analysis of State-of-the-Art Provenance-Based Intrusion Detection Systems*
  - + some other unpublished results

# What is this talk about?

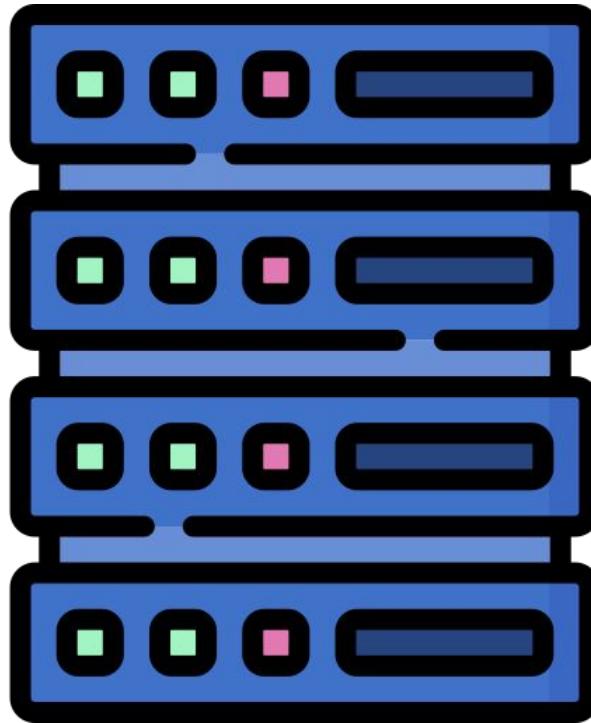
Mainly on 2 papers published at **USENIX Sec'25**.

- *ORTHRUS: Achieving High Quality of Attribution in Provenance-based Intrusion Detection Systems*
- *Sometimes Simpler is Better: A Comprehensive Analysis of State-of-the-Art Provenance-Based Intrusion Detection Systems*
  - + some other unpublished results

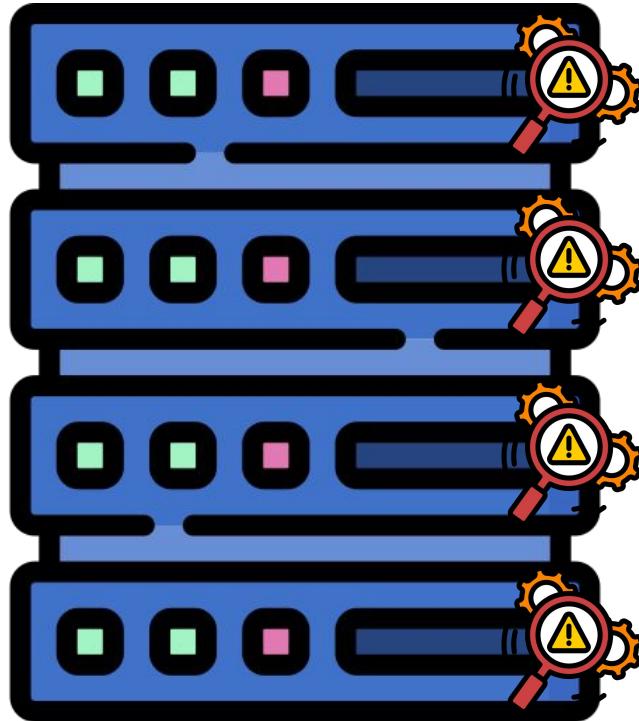
Do not hesitate to interrupt if you have a question

# Motivation

# Host-based Intrusion Detection

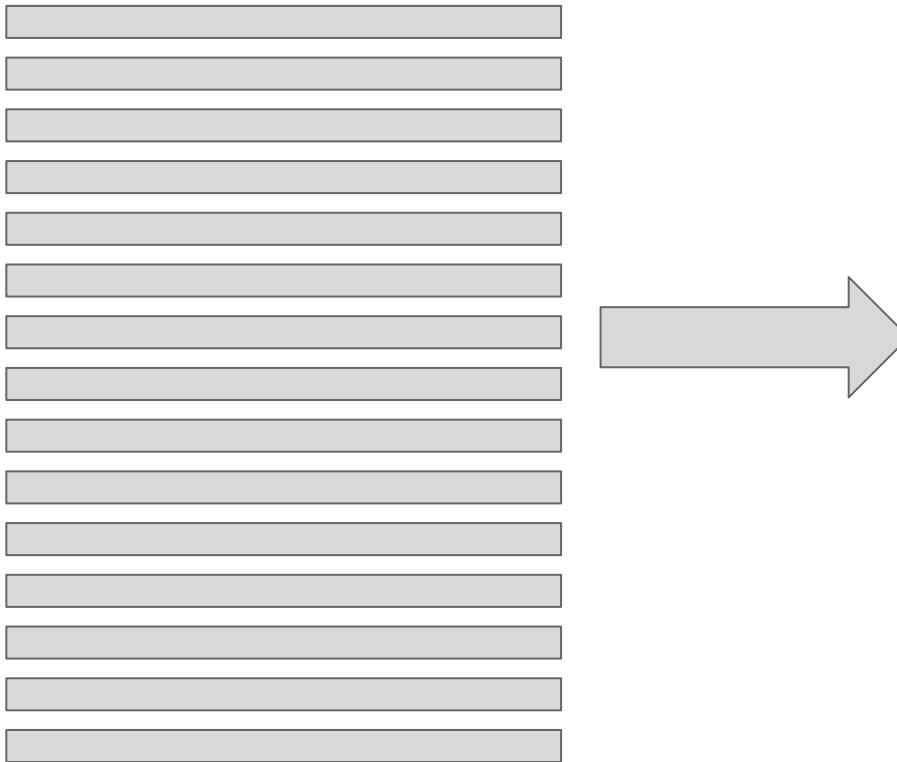


# Host-based Intrusion Detection Systems



- Deployed on individual machines
- Analyze system traces
- Look for signs of malicious behavior

# Provenance-based Intrusion Detection Systems (PIDS)



# How to build a provenance graph?

httpd receive packet from 63.169.38.150

httpd read config.ini

httpd read index.html

httpd send packet to http://63.169.38.150/

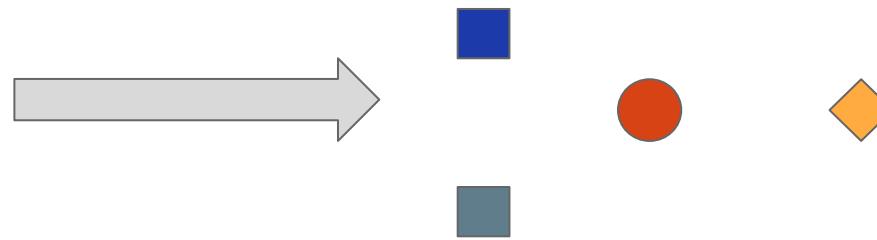
# How to build a provenance graph?

**httpd** receive packet from **63.169.38.150**

**httpd** read **config.ini**

**httpd** read **index.html**

**httpd** send packet to **63.169.38.150**



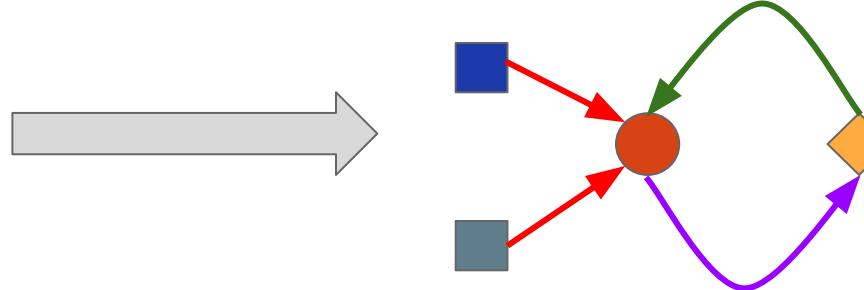
# How to build a provenance graph?

httpd receive packet from 63.169.38.150

httpd read config.ini

httpd read index.html

httpd send packet to 63.169.38.150



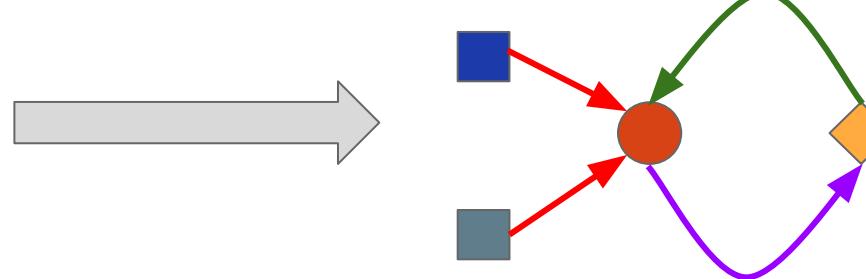
# How to build a provenance graph?

httpd receive packet from 63.169.38.150

httpd read config.ini

httpd read index.html

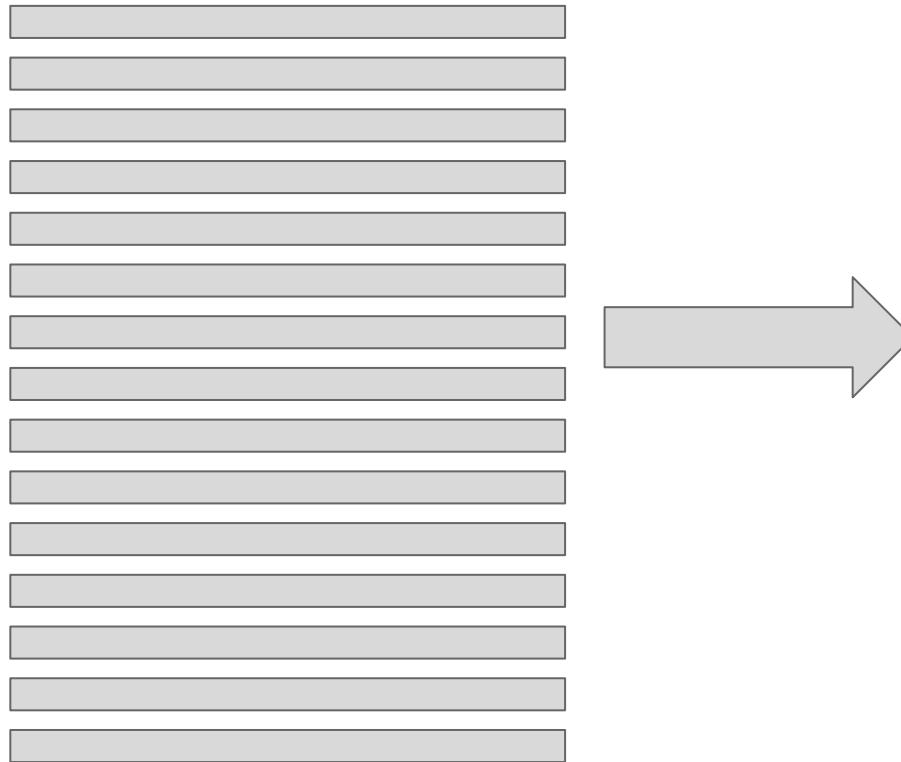
httpd send packet to 63.169.38.150



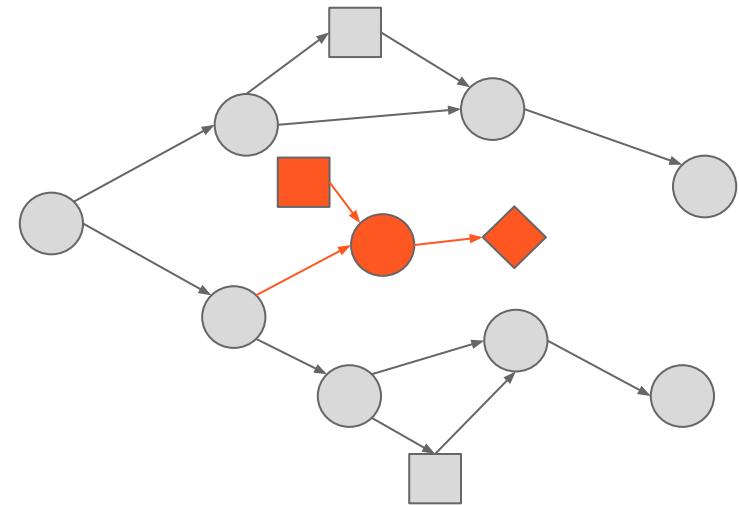
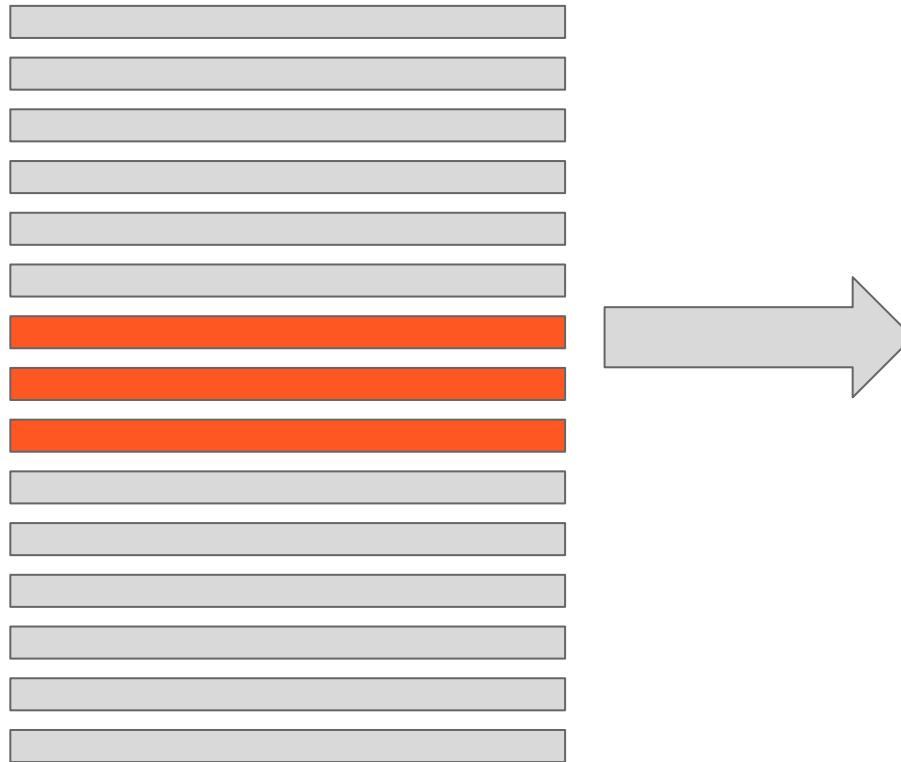
Compiling the Linux Kernel: ~2M graph elements

**What is the  
intuition behind  
PIDS?**

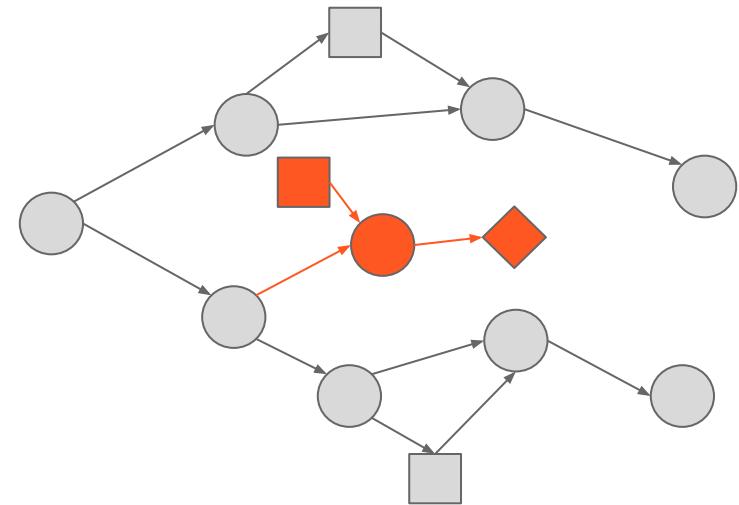
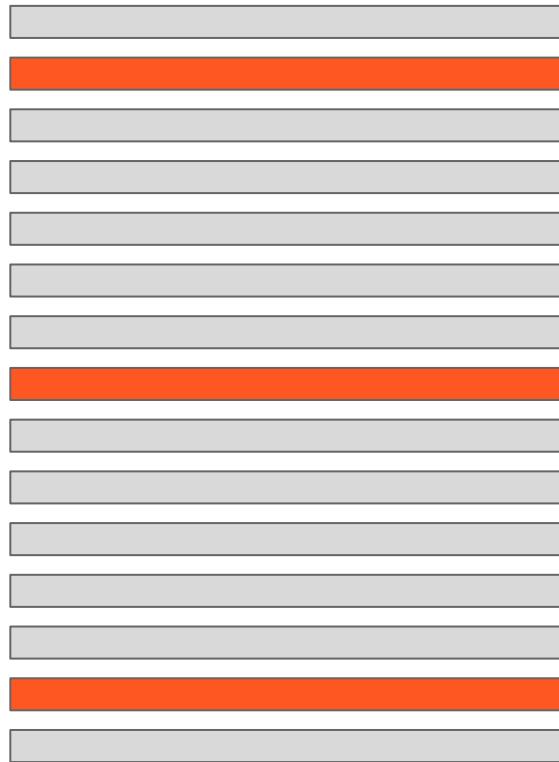
# What is the intuition behind PIDS?



# What is the intuition behind PIDS?

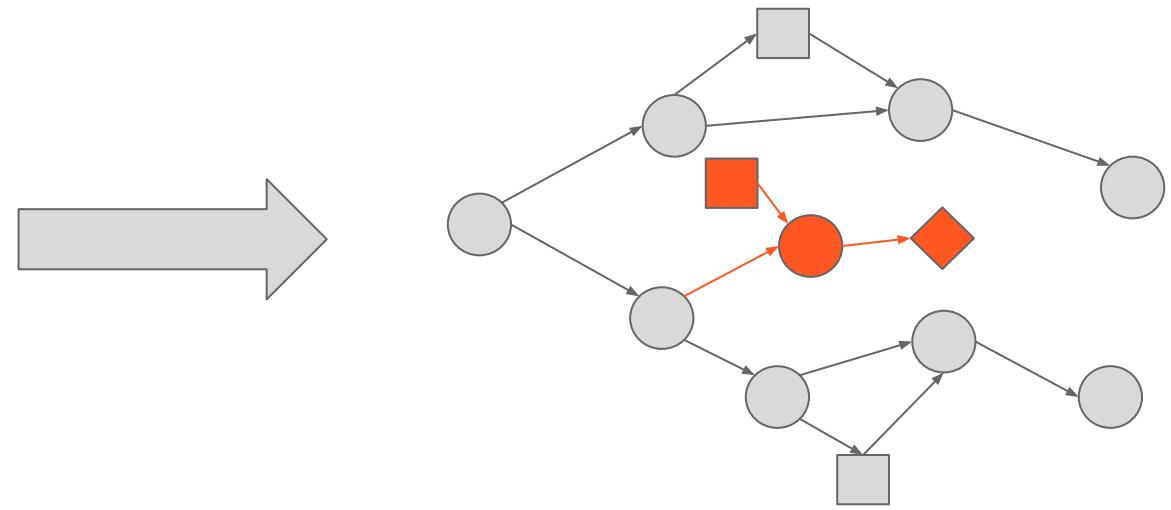
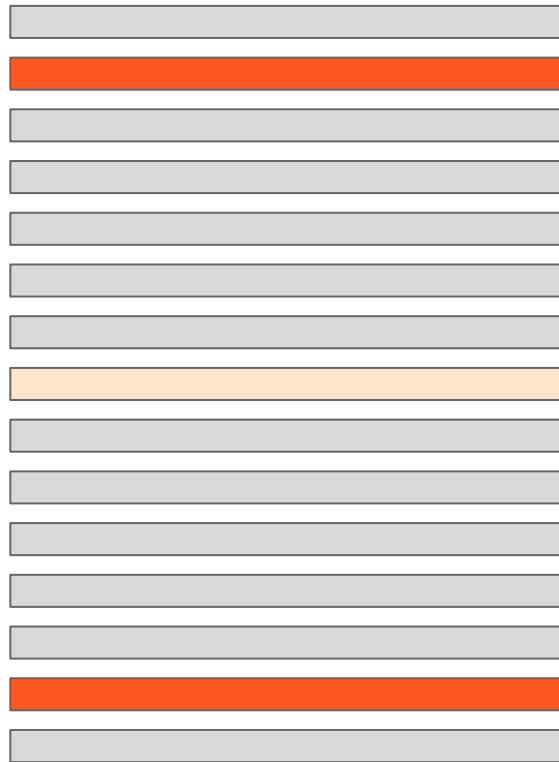


# What is the intuition behind PIDS?



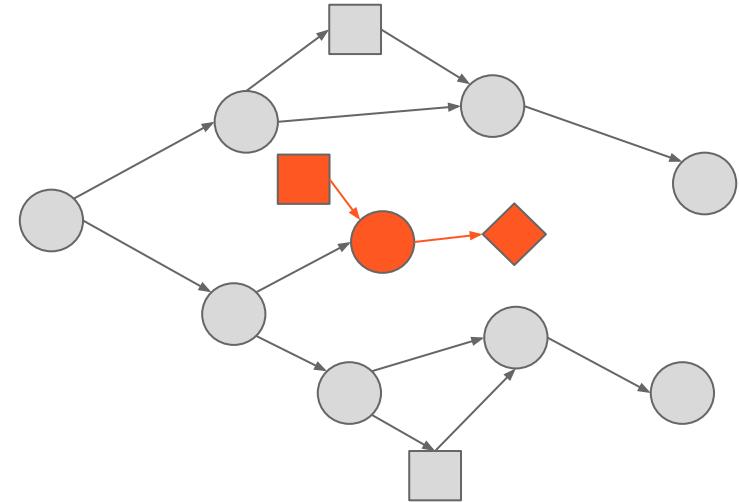
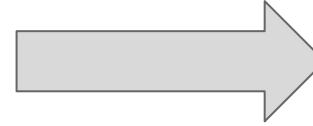
Hide with benign events

# What is the intuition behind PIDS?



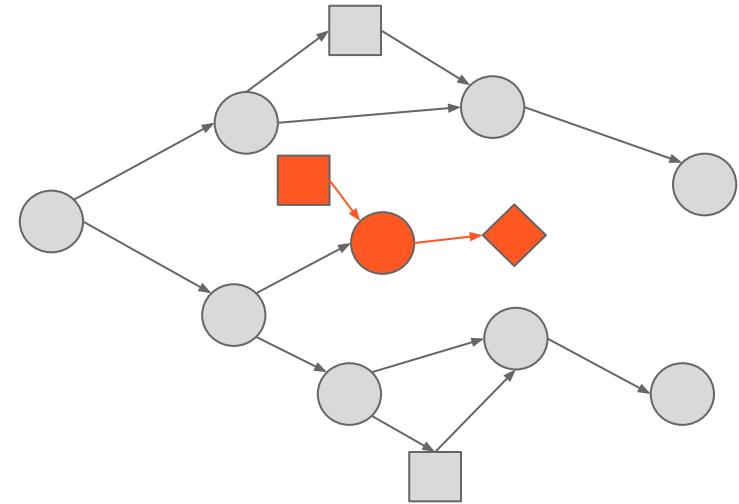
Masquerading as a benign event

# What is the intuition behind PIDS?



Attacks spaced in time

# What is the intuition behind PIDS?



Causality relationship is preserved.

# Two potential approaches

## Signature-based detection

Match known malicious graph patterns

## Anomaly-based detection

Detect pattern that deviate from normal behavior

+ Higher precision

+ Can detect unknown patterns

- Only detect known pattern

- Lower precision

# Two potential approaches

## Signature-based detection

Match known malicious graph patterns

- + Higher precision

- Only detect known pattern

## Anomaly-based detection

Detect patterns that deviate from normal behavior

- + Can detect unknown patterns

- Lower precision

PIDS emerged to detect APT, zero-day exploit, etc.

# Anomaly-based detection

- Train on past benign activity
  - Self-supervised learning
- During inference, perform anomaly detection
  - E.g., high reconstruction loss

# Limitation of SOTA methods

- Security analysts are overwhelmed with false positives
  - Alert Fatigue - Wajih et al., NDSS 2019
  - Burn out - Chandran et al., SOUPS 2015



# Our goal

- Security analysts are overwhelmed with false positives
  - Alert Fatigue - Wajih et al., NDSS 2019
  - Burn out - Chandran et al., SOUPS 2015



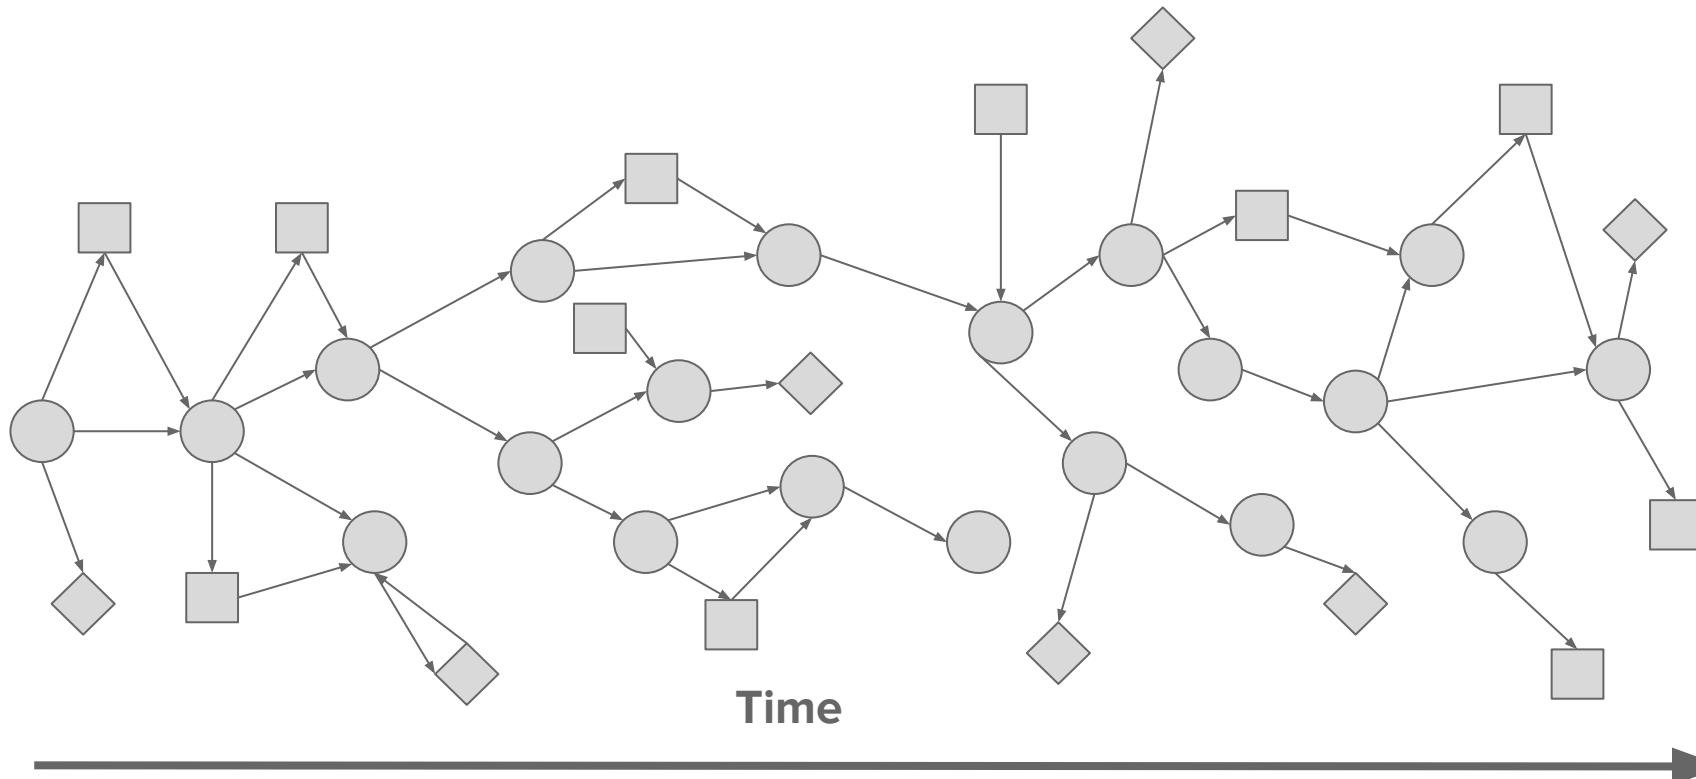
Reduce the amount of information they need to go through

# **Understanding Attribution Quality**

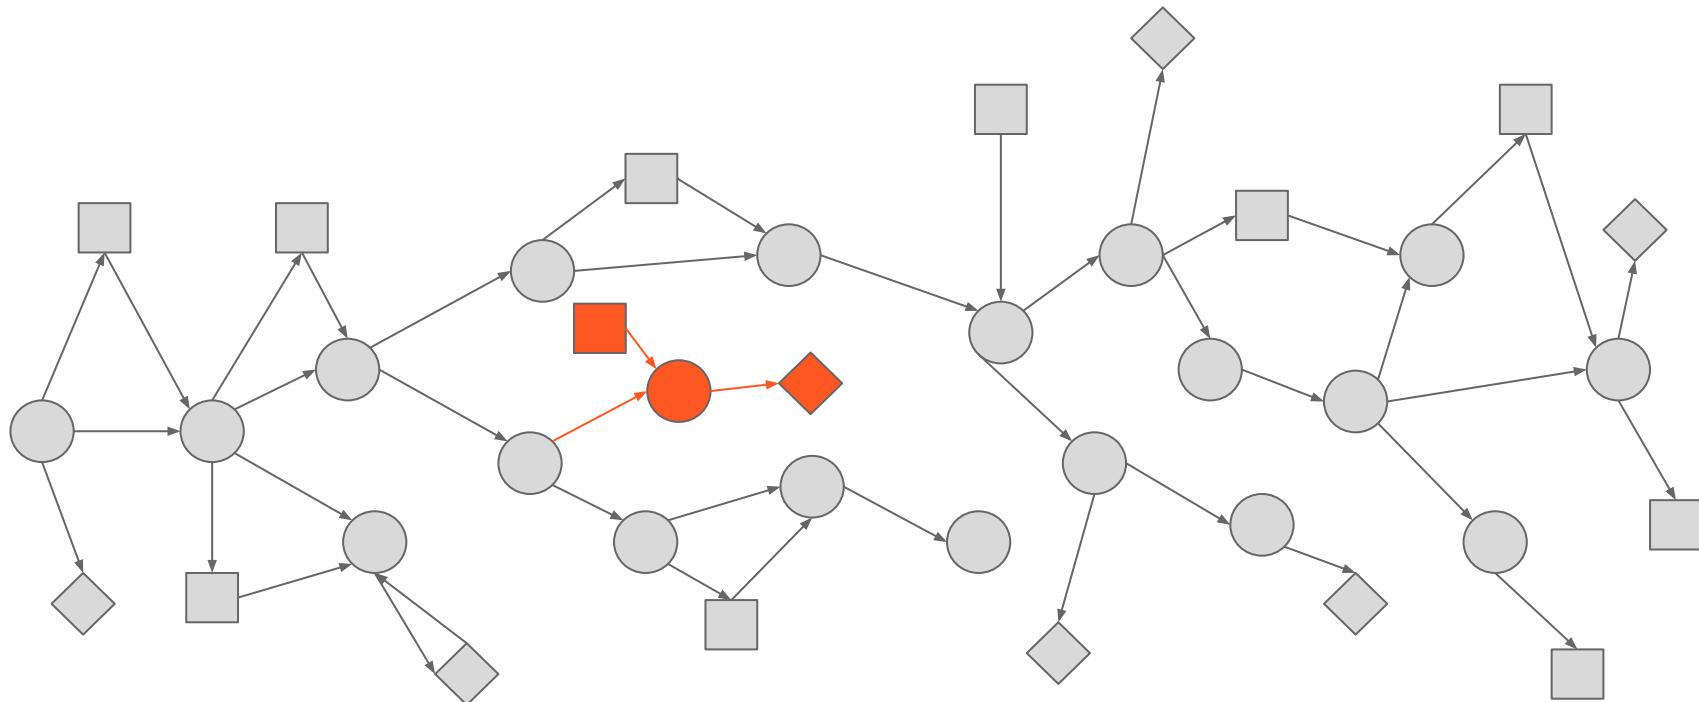
# Definition

*Attribution Quality* refers to the **amount of effort** required from a human analyst to investigate an IDS's detection output, uncover the root causes of an attack, and dismiss potential false alarms.

# Assume a provenance graph

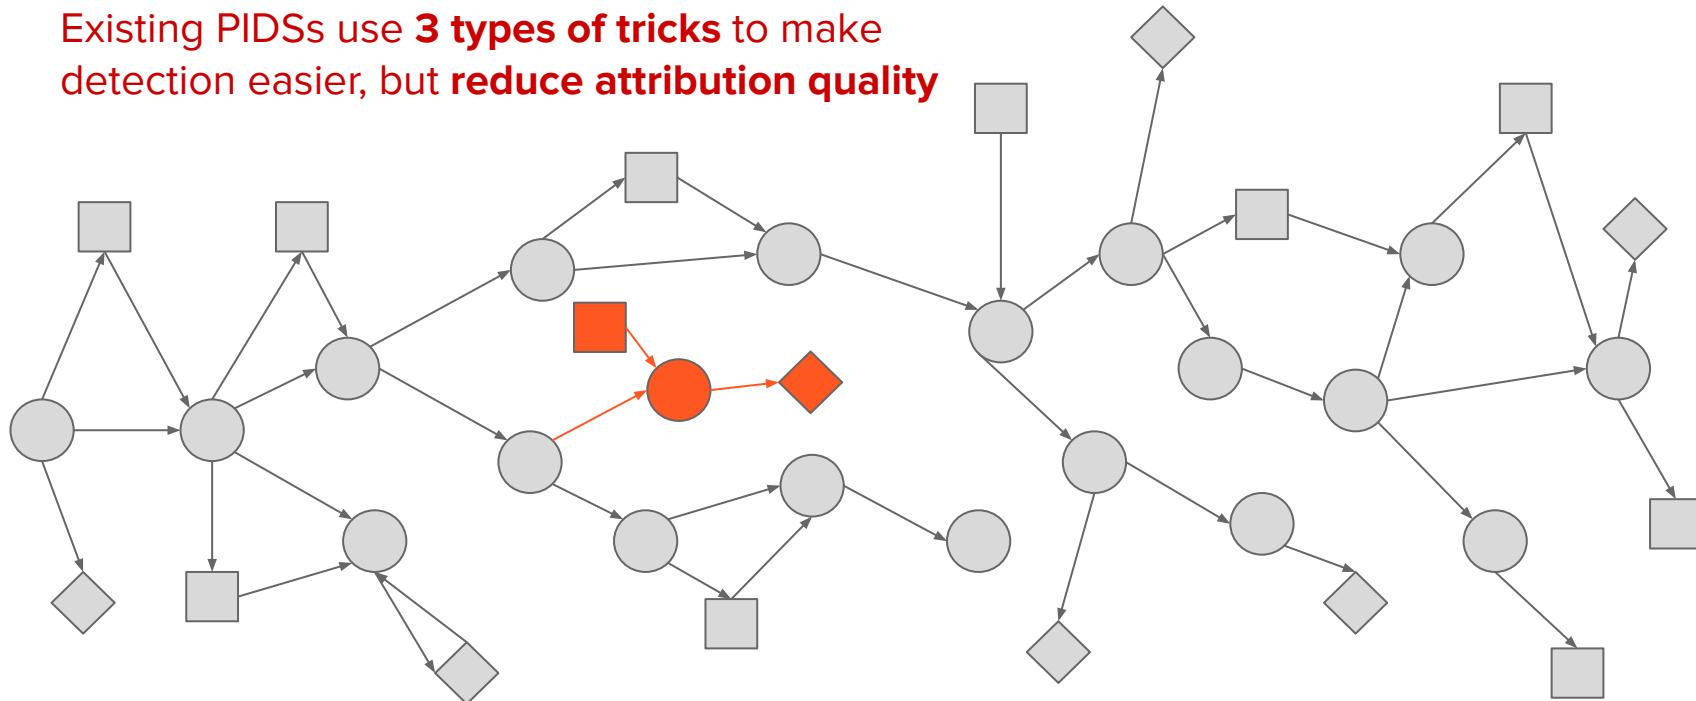


# Assume an attack



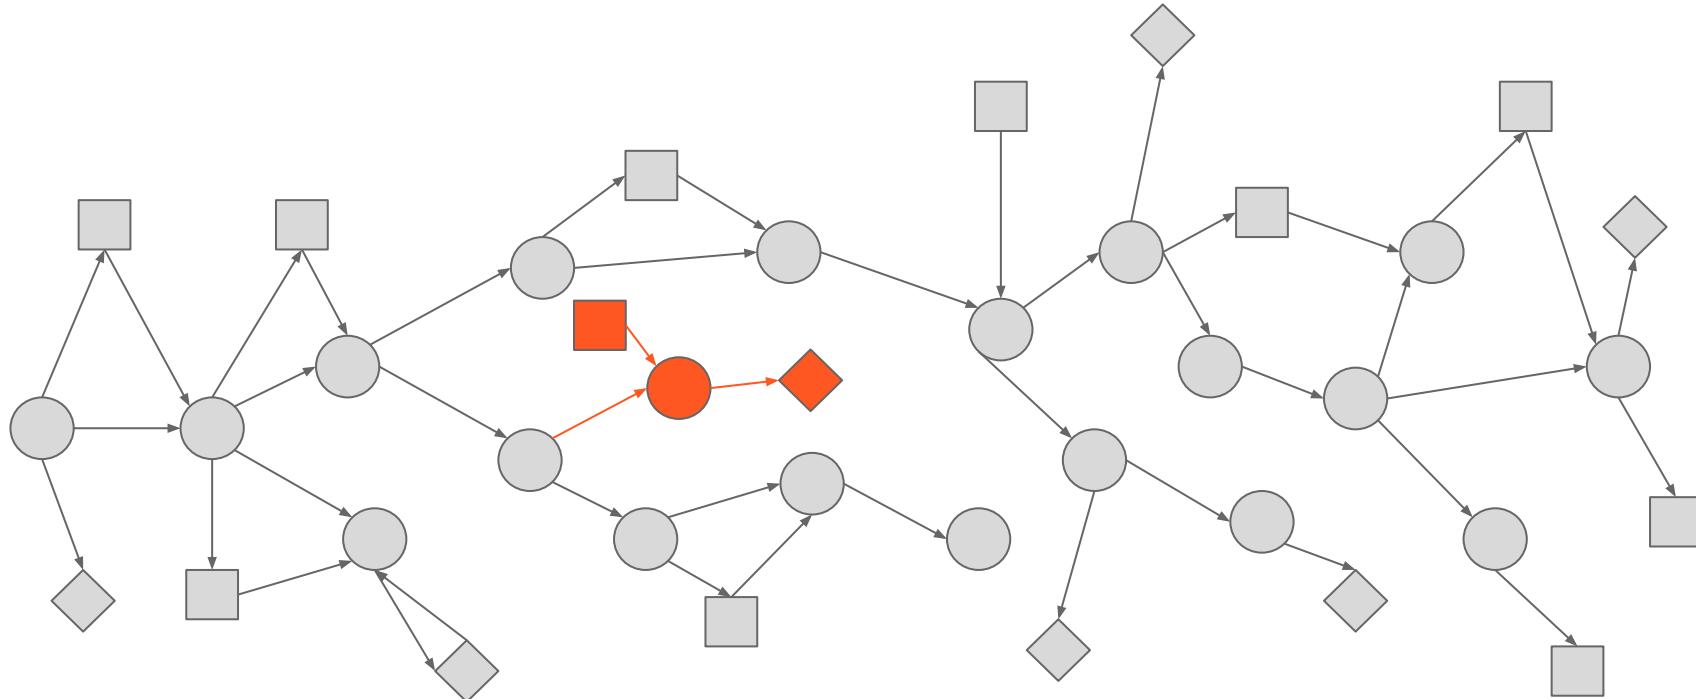
# Assume an attack

Existing PIDSs use **3 types of tricks** to make detection easier, but **reduce attribution quality**



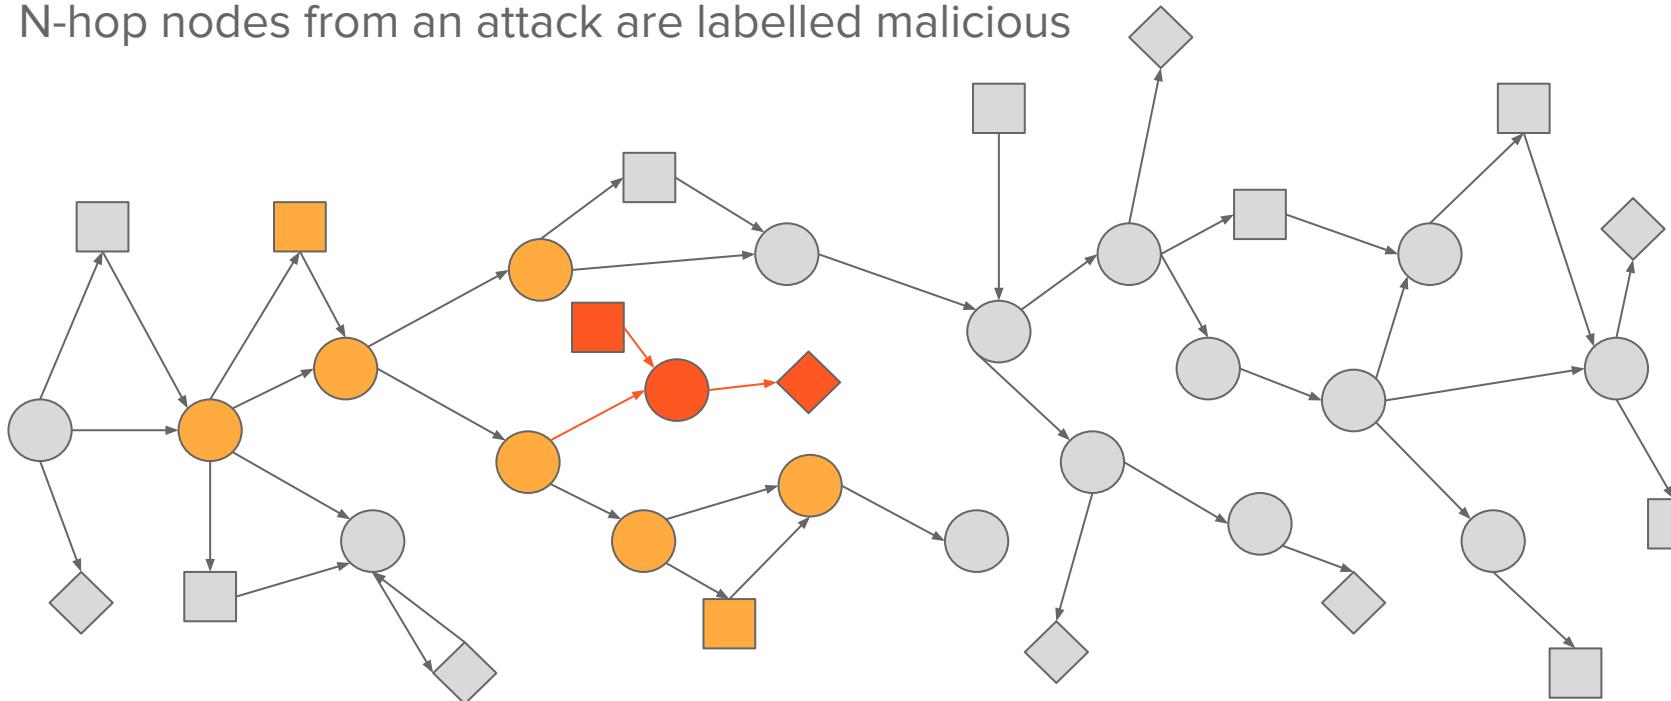
# 1. Neighborhood Attribution

Concerned PIDSs:  
**Flash**, S&P 2024  
**MAGIC**, USENIX Sec 2024  
**ThreaTrace**, TIFS 2022



# 1. Neighborhood Attribution

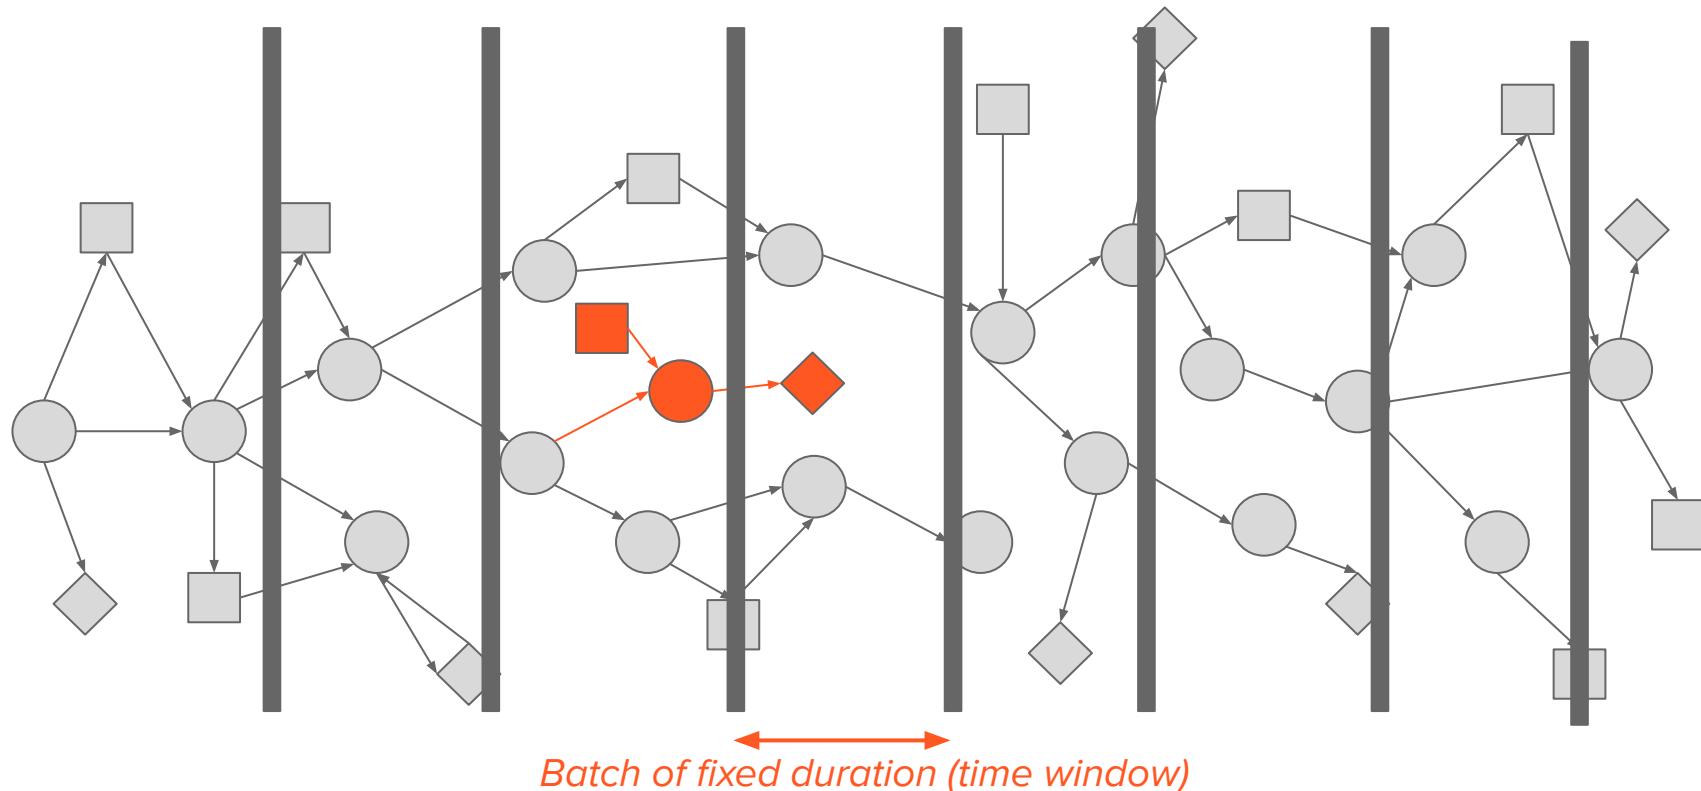
N-hop nodes from an attack are labelled malicious



Concerned PIDSs:  
**Flash**, S&P 2024  
**MAGIC**, USENIX Sec 2024  
**ThreaTrace**, TIFS 2022

## 2. Batch Attribution

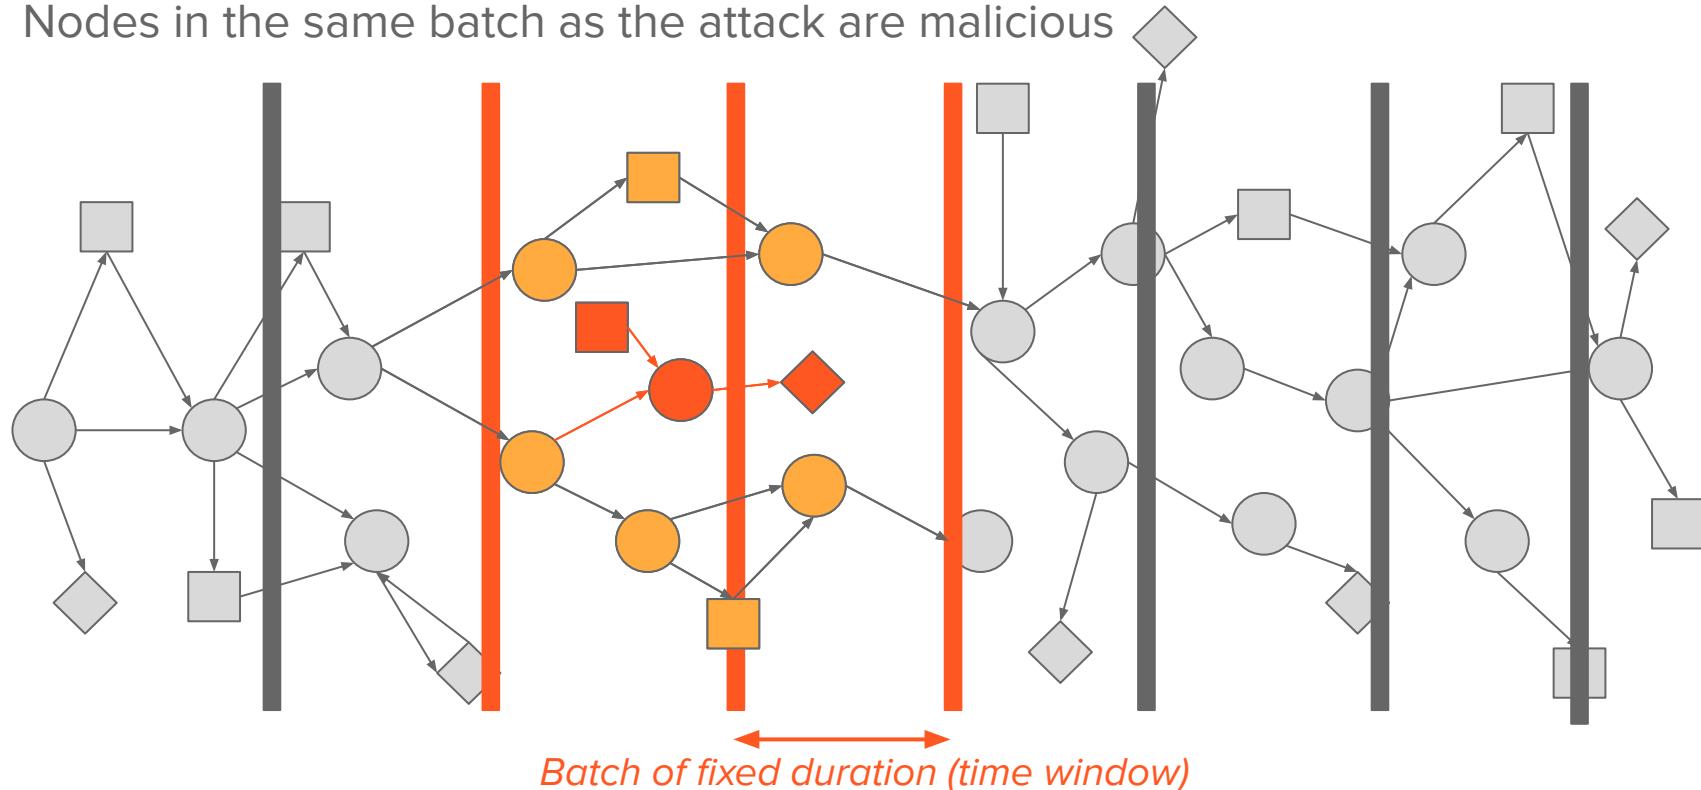
Concerned PIDSs:  
**Kairos**, S&P 2024  
**EdgeTorrent**, RAID's 2023



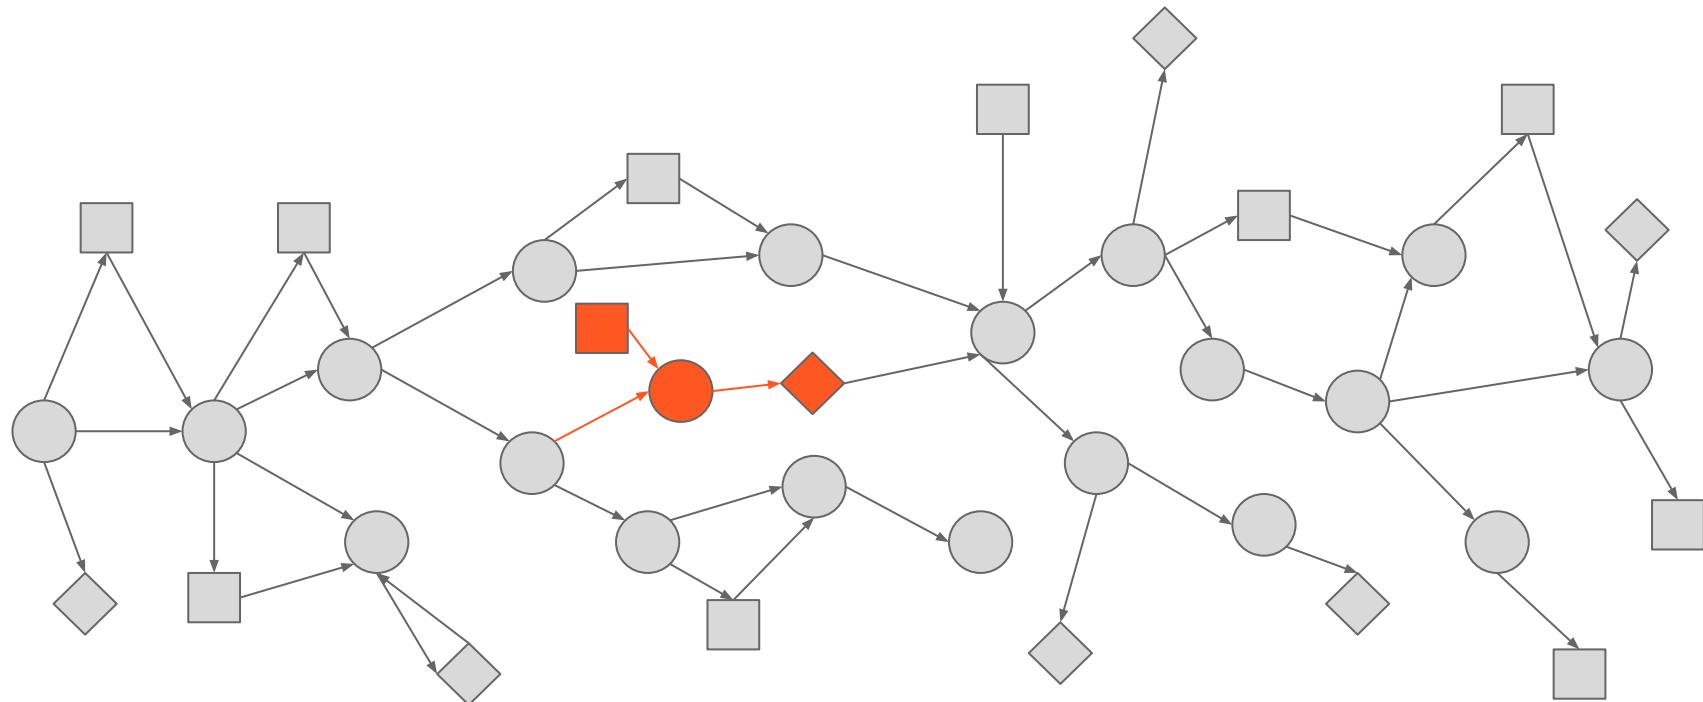
## 2. Batch Attribution

Concerned PIDSs:  
**Kairos**, S&P 2024  
**EdgeTorrent**, RAID's 2023

Nodes in the same batch as the attack are malicious

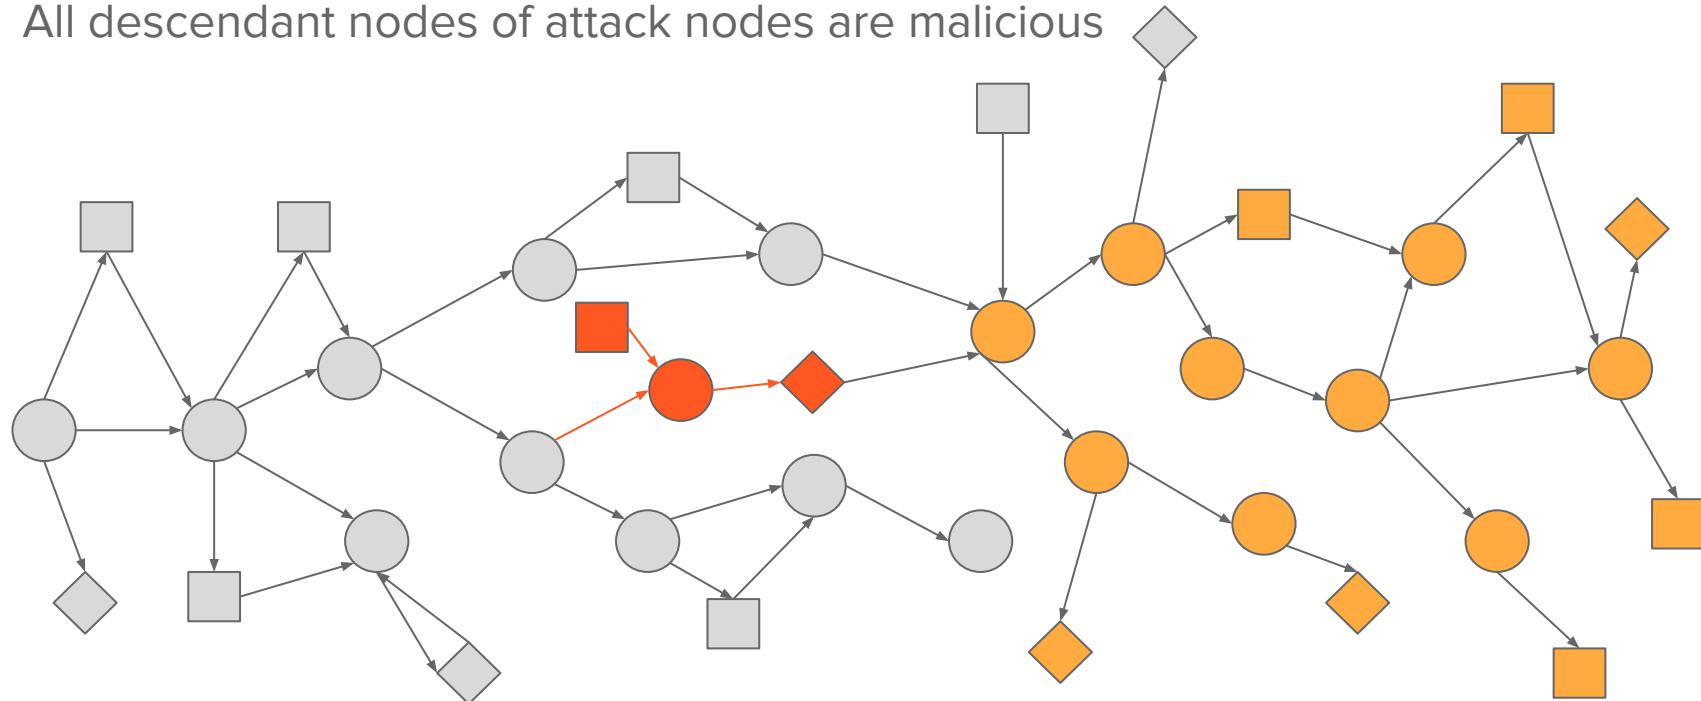


### 3. Source Attribution



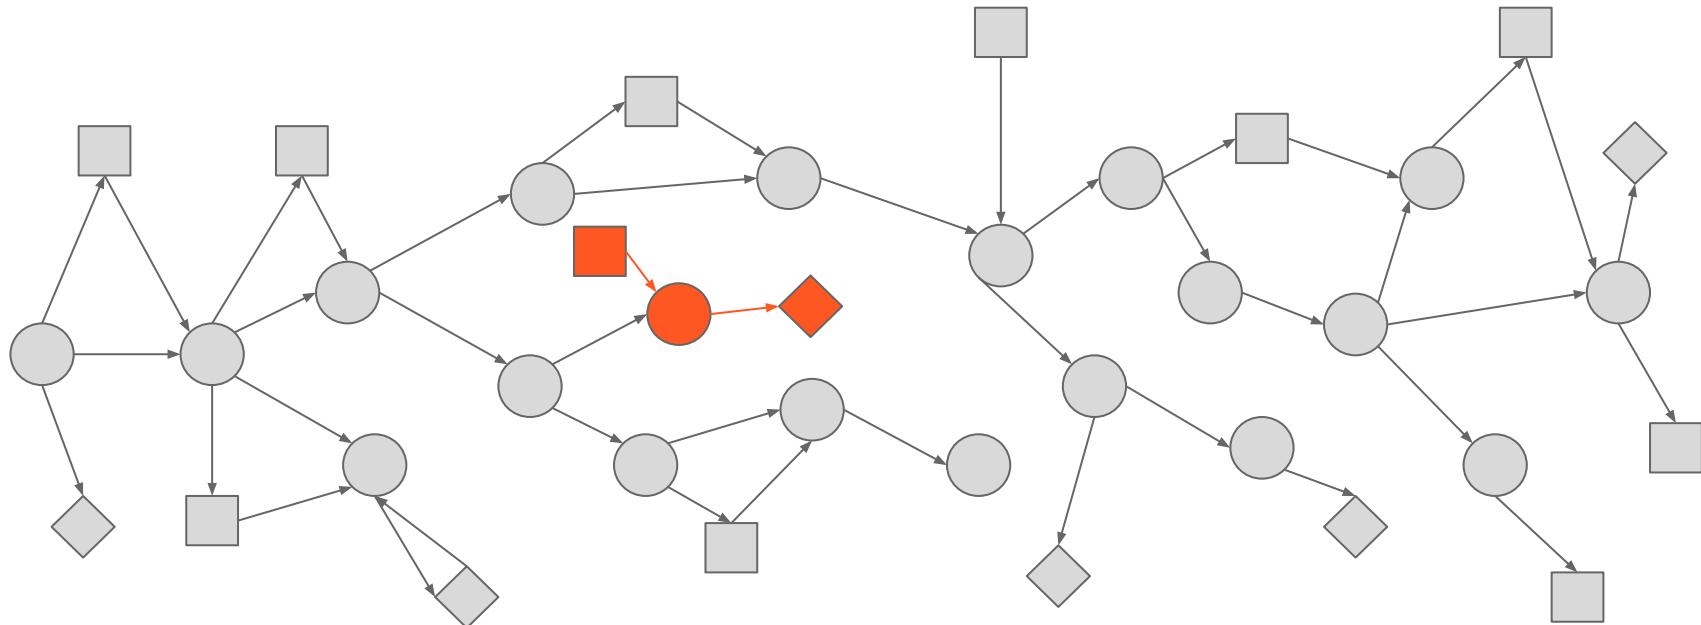
### 3. Source Attribution

All descendant nodes of attack nodes are malicious



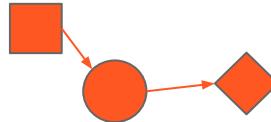
# Our Strategy

No tricks, we directly use the attack nodes from dataset ground truths



# Our Strategy

No tricks, we directly use the attack nodes from dataset ground truths



- Much harder detection (few nodes)
- Concise detection reports (less work for analysts)

# Idealized Detection Performance

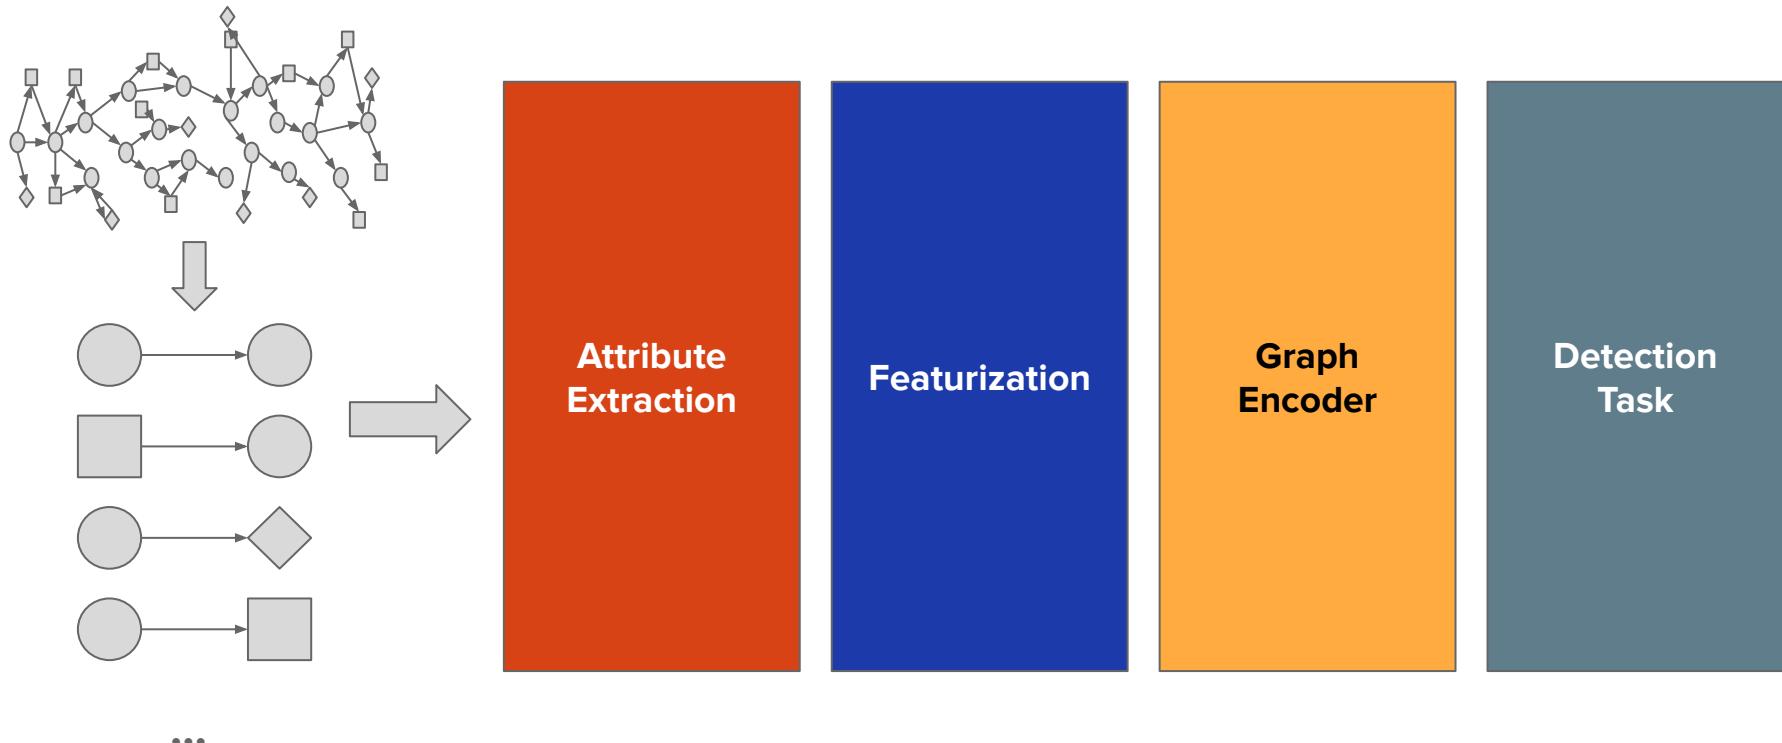
- Assuming perfect detection based on their design
- Past systems overwhelm security analyst with large alerts

Number of attack nodes to detect per attribution strategy:

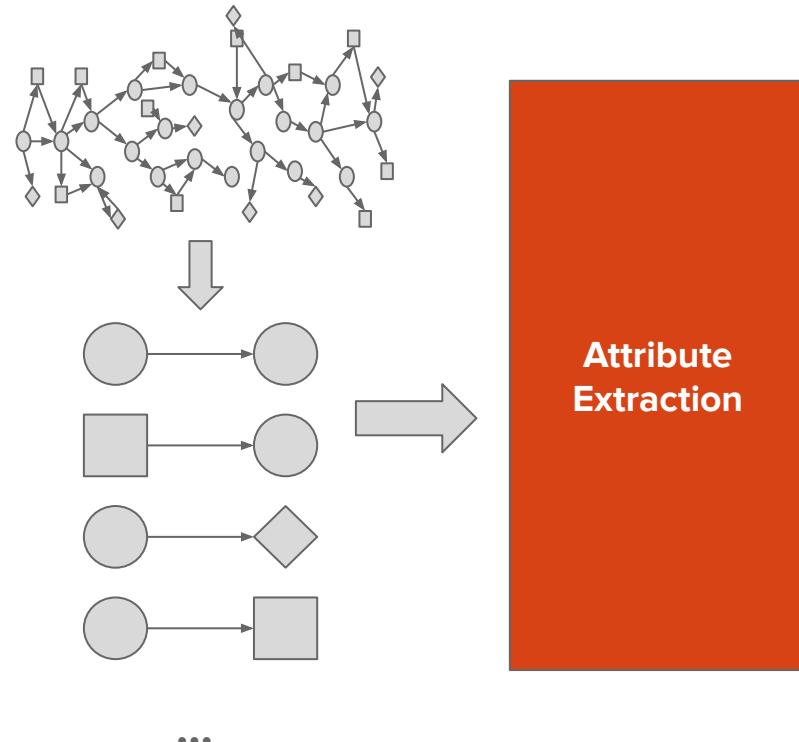
Dataset	Total Nodes	Neighborhood	Batch	Source	Node-level (Ours)
E5-CADETS	7,632,792	20,524	717,783	401,065	123
E5-THEIA	1,728,121	162,714	61,368	9,374	69
E5-CLEARSCOPE	326,338	48,488	8,636	1,020	51

# **Orthrus: a PIDS for node-level detection**

# Orthrus design



# Orthrus design

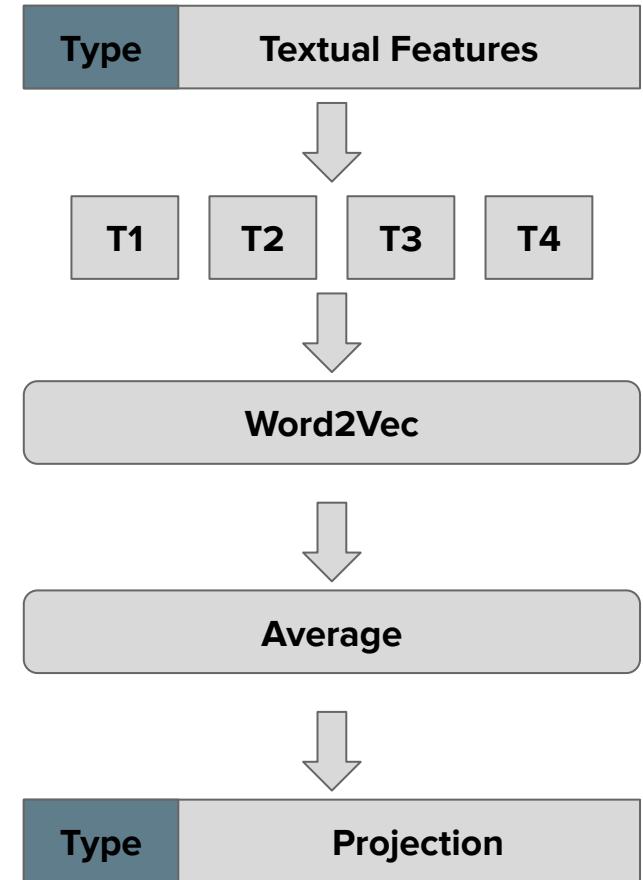
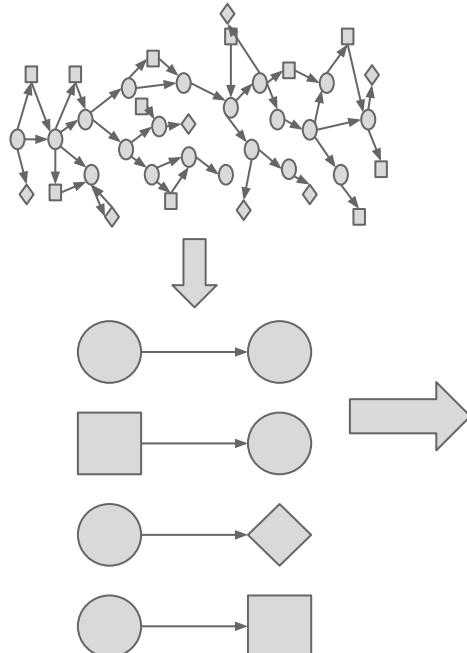


**File:** type + path (e.g. /etc/passwd)

**Process:** type + cmd line (e.g. ls -l -t -r)

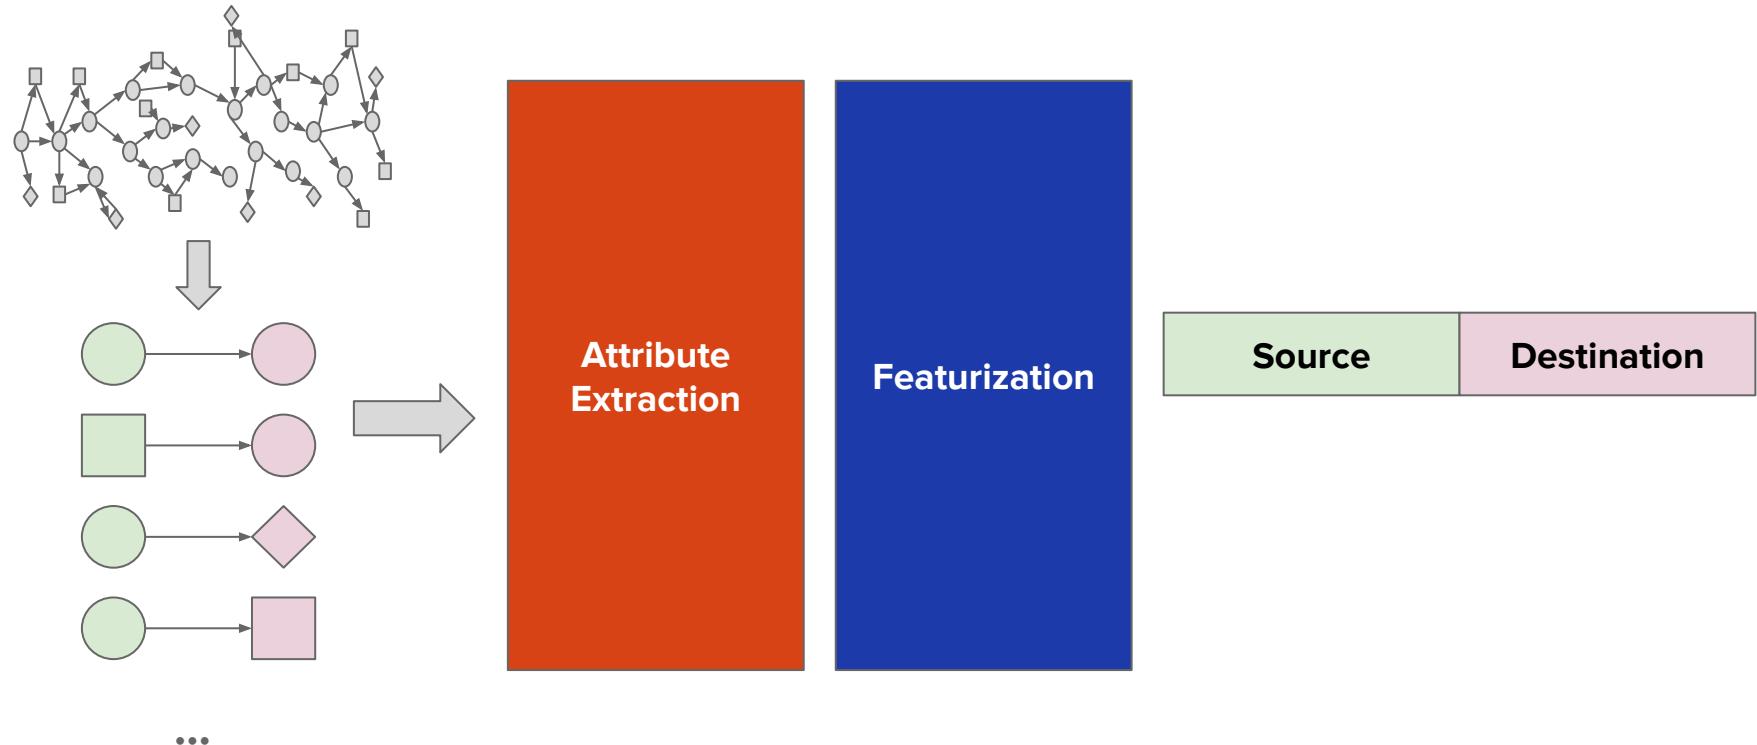
**Netflow:** type + IP + port (e.g. 192.168.1.2 80)

# Orthrus design

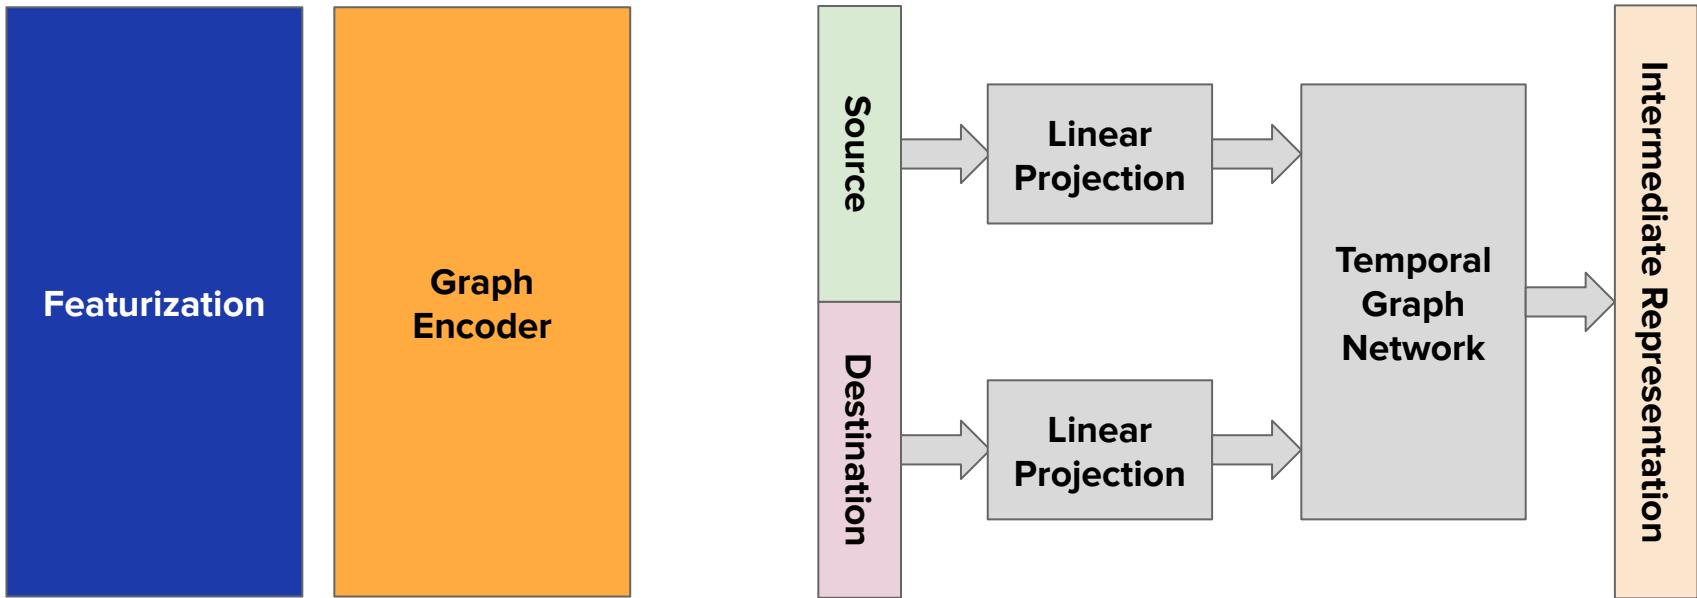


...

# Orthrus design

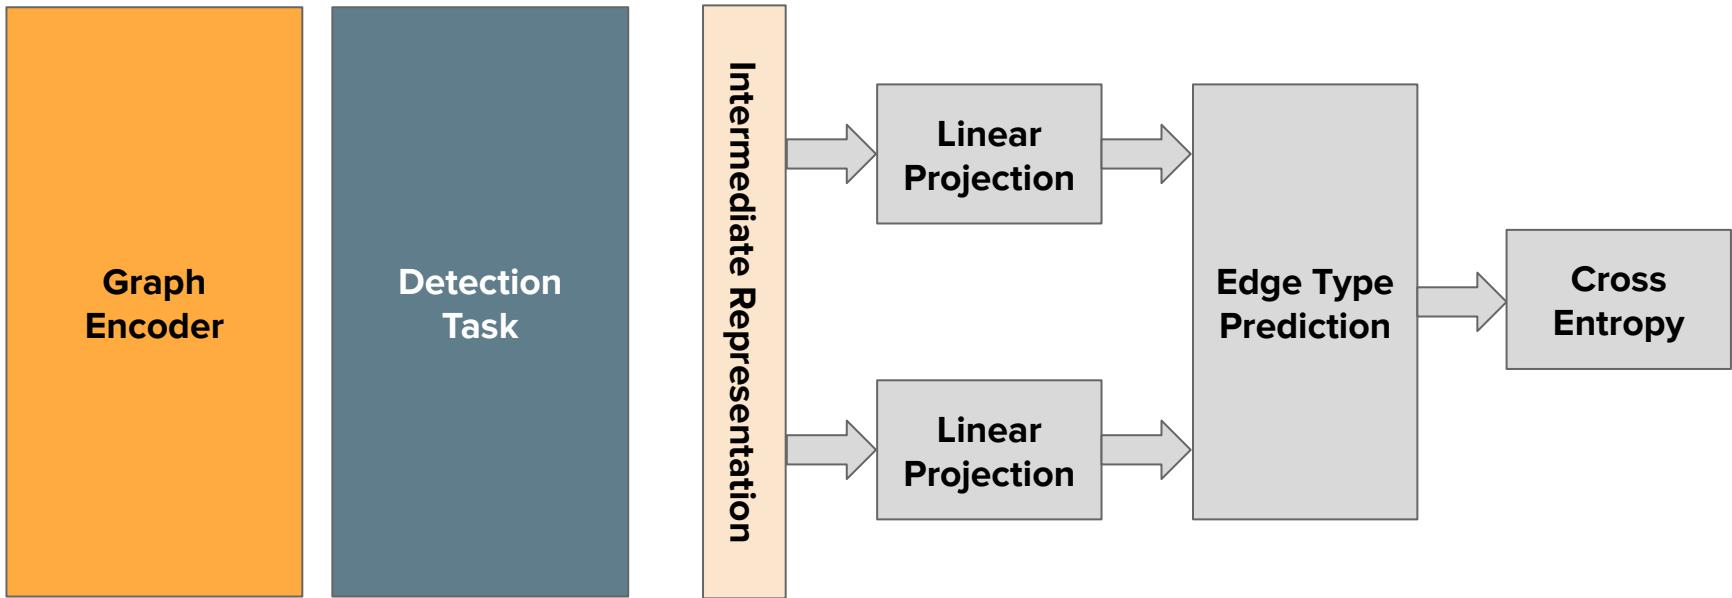


# Orthrus design



We made some changes to TGN to fit our problem.

# Orthrus design



# Evaluation

# Baselines

- We picked **5** graph learning-based **SOTA PIDSs**
- Hyperparameters tuned consistently
- Experiments based on our node-level labels

# Results

## E3 Datasets

System	E3-CADETS (TP / FP / Prec)	E3-THEIA	E3-CLEARSCOPE
Kairos	0 / 9 / 0.00	4 / 0 / 1.00	0 / 7 / 0.00
Threatrace	61 / 252k / 0.00	88 / 672k / 0.00	41 / 88k / 0.00
SIGL	0 / 80 / 0.00	1 / 29 / 0.03	1 / 11k / 0.00
MAGIC	63 / 80k / 0.00	115 / 395k / 0.00	40 / 102k / 0.00
Flash	13 / 2.4k / 0.01	22 / 32k / 0.00	0 / 15k / 0.00
Orthrus	<b>10 / 0 / 1.00</b>	<b>8 / 0 / 1.00</b>	<b>1 / 1 / 0.50</b>

## E5 Datasets

System	E5-CADETS (TP / FP / Prec)	E5-THEIA	E5-CLEARSCOPE
Kairos	0 / 6 / 0.00	0 / 2 / 0.00	1 / 7 / 0.25
Threatrace	91 / 3M / 0.00	66 / 739k / 0.00	41 / 142k / 0.00
SIGL	0 / 66 / 0.00	0 / 23 / 0.00	10 / 63 / 0.14
MAGIC	123 / 3M / 0.00	1 / 297k / 0.00	51 / 139k / 0.00
Flash	45 / 34k / 0.00	43 / 296k / 0.00	15 / 4.6k / 0.00
Orthrus	<b>1 / 5 / 0.17</b>	<b>2 / 0 / 1.00</b>	<b>2 / 3 / 0.22</b>

# Results

- Orthrus can detect **all attacks** in each dataset
- It detects a few nodes only, but with high precision
- Attack reconstruction algorithms for provenance graphs can be used

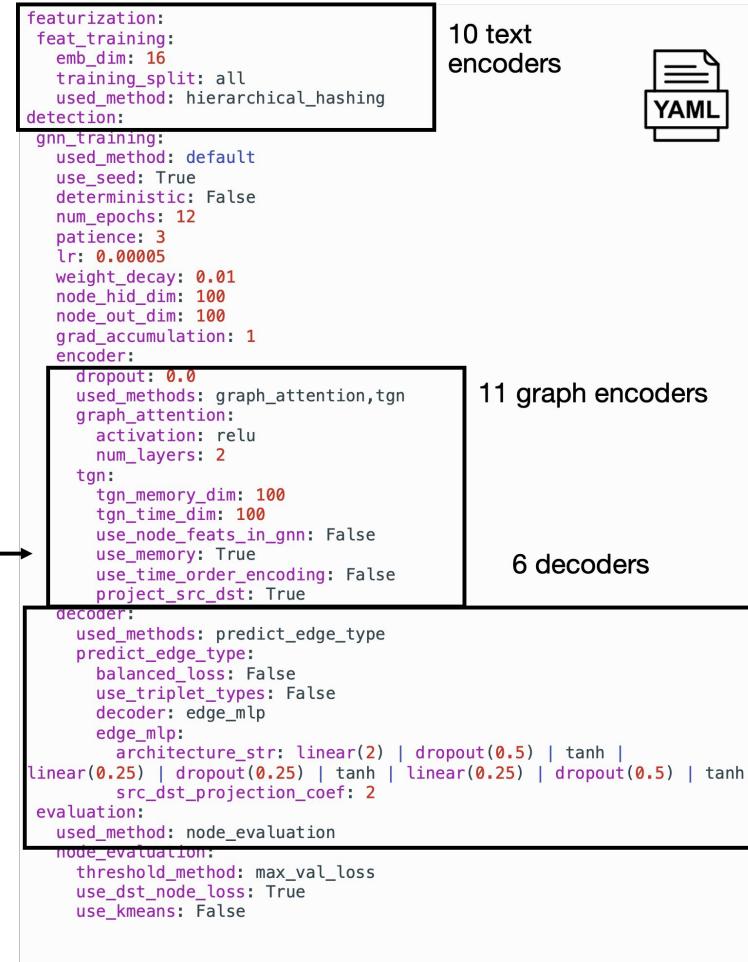
# Toward More Practical PIDSs

# PIDSMaker

## A framework to design PIDSSs

We reimplemented 7 SOTA PIDSSs in a unified (open source) framework

1. Kairos
2. ThreaTrace
3. NodLink
4. Magic
5. Flash
6. R-Caid
7. SIGL
8. Orthrus



# PIDSMaker

## A framework to design PIDSSs

### Goal:

- Combinatorial architecture search
- Consistent evaluation across papers

```
featurization:  
feat_training:  
emb_dim: 16  
training_split: all  
used_method: hierarchical_hashing  
detection:
```

10 text  
encoders



```
gnn_training:  
used_method: default  
use_seed: True  
deterministic: False  
num_epochs: 12  
patience: 3  
lr: 0.00005  
weight_decay: 0.01  
node_hid_dim: 100  
node_out_dim: 100  
grad_accumulation: 1  
encoder:
```

```
dropout: 0.0  
used_methods: graph_attention,tgn  
graph_attention:  
activation: relu  
num_layers: 2  
tgn:  
tgn_memory_dim: 100  
tgn_time_dim: 100  
use_node_feats_in_gnn: False  
use_memory: True  
use_time_order_encoding: False  
project_src_dst: True
```

11 graph encoders

```
decoder:  
used_methods: predict_edge_type  
predict_edge_type:  
balanced_loss: False  
use_triplet_types: False  
decoder: edge_mlp  
edge_mlp:  
architecture_str: linear(2) | dropout(0.5) | tanh |  
linear(0.25) | dropout(0.25) | tanh | linear(0.25) | dropout(0.5) | tanh  
src_dst_projection_coef: 2  
evaluation:  
used_method: node_evaluation  
node_evaluation:  
threshold_method: max_val_loss  
use_dst_node_loss: True  
use_kmeans: False
```

6 decoders

# Addressing Existing Shortcomings

- We found 9 key shortcomings that hinder practicality of SOTA PIDSs

System	SC <sub>1</sub>	SC <sub>2</sub>	SC <sub>3</sub>	SC <sub>4</sub>	SC <sub>5</sub>	SC <sub>6</sub>	SC <sub>7</sub>	SC <sub>8</sub>	SC <sub>9</sub>
SIGL			✓			✓			
THREATTRACE						✓			
NODLINK	✓		✓			✓			
MAGIC						✓			
KAIROS						✓			
FLASH						✓			
R-CAID			✓			✓			
ORTHRUS	✓	✓							

# Addressing Existing Shortcomings

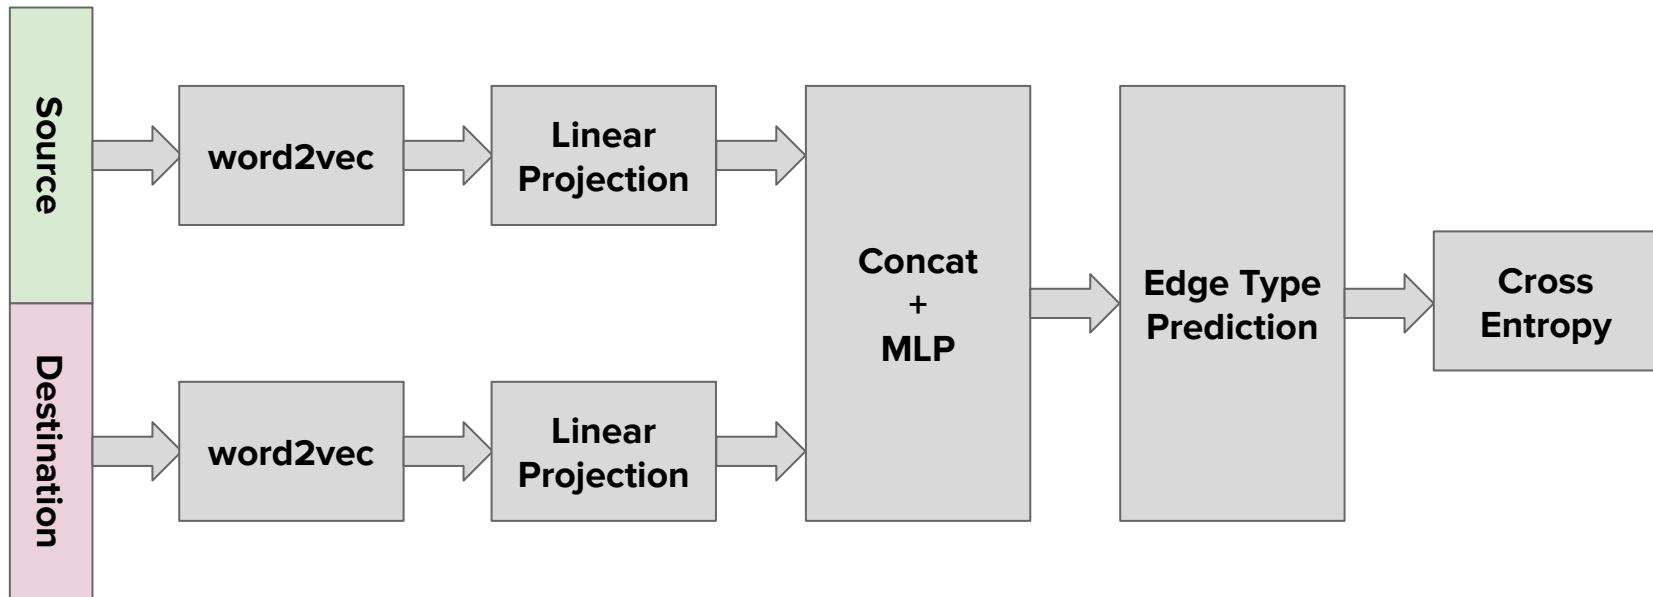
- We found 9 key shortcomings that hinder practicality of SOTA PIDSs

System	SC <sub>1</sub>	SC <sub>2</sub>	SC <sub>3</sub>	SC <sub>4</sub>	SC <sub>5</sub>	SC <sub>6</sub>	SC <sub>7</sub>	SC <sub>8</sub>	SC <sub>9</sub>
SIGL			✓			✓			
THREATTRACE						✓			
NODLINK	✓		✓			✓			
MAGIC						✓			
KAIROS						✓			
FLASH						✓			
R-CAID				✓		✓			
ORTHRUS	✓	✓							
VELOX	✓	✓	✓	✓	✓	✓	✓	✓	✓



We designed **Velox**, a PIDS that addresses all these shortcomings

# Velox design

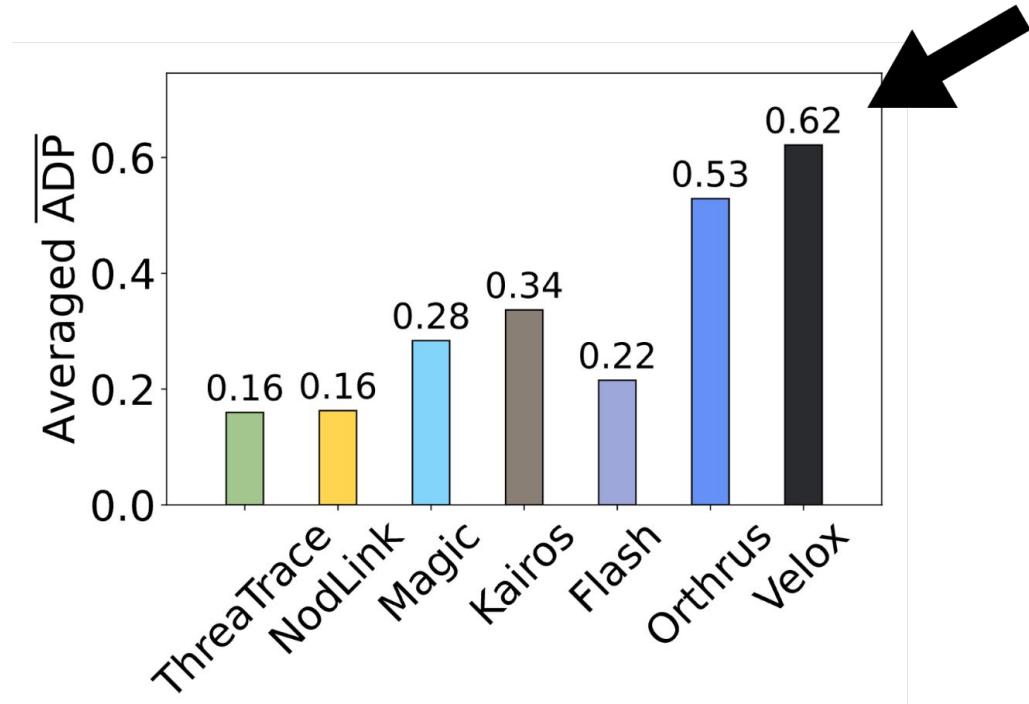


- This design has been selected after **~453 days of GPU compute**
- Unlike all other PIDSs, no complex graph encoder is used

# An Unexpected Discovery

Velox has the **highest** average  
**Attack Detection Precision**  
**(ADP)** across 9 DARPA datasets

**ADP:** measures the ability to  
detect all attacks in a dataset  
with high precision



*ADP averaged across 9 DARPA datasets*

# Discussion

- Complex architectures are not always needed
- APTs can be detected using textual features features only!
  - Realistic attacks?

# **Future Directions**

# Limitations & Future Work

Limitation	Future Work
• Benign anomalies could be false positives	→ Design datasets with benign anomalies
• Training instability	→ Add some supervision during training

# Want to play with PIDSMaker?

- PIDSMaker is open source
- 8 PIDSS (incl. Orthrus & Velox)
- 9 pre-processed datasets
- You can build your own models

<https://github.com/ubc-provenance/PIDSMaker>

The first framework designed to build and experiment with provenance-based intrusion detection systems (PIDSS) using deep learning architectures. It provides a single codebase to run most recent state-of-the-arts systems and easily customize them to develop new variants.

Currently supported PIDSS:

- **Velox** (USENIX Sec'25): [Sometimes Simpler is Better: A Comprehensive Analysis of State-of-the-Art Provenance-Based Intrusion Detection Systems](#)
- **Orthrus** (USENIX Sec'25): [ORTHRUS: Achieving High Quality of Attribution in Provenance-based Intrusion Detection Systems](#)
- **R-Caid** (IEEE S&P'24): [R-CAID: Embedding Root Cause Analysis within Provenance-based Intrusion Detection](#)
- **Flash** (IEEE S&P'24): [Flash: A Comprehensive Approach to Intrusion Detection via Provenance Graph Representation Learning](#)
- **Kairos** (IEEE S&P'24): [Kairos: Practical Intrusion Detection and Investigation using Whole-system Provenance](#)
- **Magic** (USENIX Sec'24): [MAGIC: Detecting Advanced Persistent Threats via Masked Graph Representation Learning](#)
- **NodLink** (NDSS'24): [NODLINK: An Online System for Fine-Grained APT Attack Detection and Investigation](#)
- **ThreaTrace** (IEEE TIFS'22): [THREATTRACE: Detecting and Tracing Host-Based Threats in Node Level Through Provenance Graph Learning](#)

# Collaborations

## Want to work with us?

We are looking for motivated students to collaborate!

- Fully-funded PhD offers are still available in our team (UBC x Amazon)
- Possible to join as a visiting research student too

Drop me an e-mail: [tristan.bilot@universite-paris-saclay.fr](mailto:tristan.bilot@universite-paris-saclay.fr)